

Architecture Impact on Performance

18-645: How to write fast code

Tze Meng Low

© 2025 Tze Meng Low

1

SOFTWARE

18645 – How to write fast code

HARDWARE

© 2025 Tze Meng Low

2

What is software...

$$\boxed{C} = \boxed{A} \times \boxed{B}$$

C/C++

```
for (int i = 0; i != m; ++i)
  for (int j = 0; j != n; ++j)
    for (int k = 0; k != p; ++k)
      c[j*m+i] += a[p*m+i] * b[j*p+k];
```

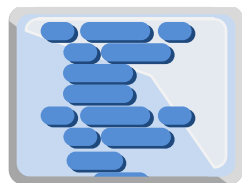
Python

```
C = A @ B
```

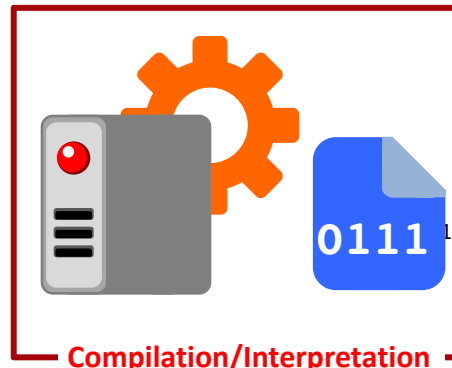
© 2025 Tze Meng Low

3

Software wrote is not software executed



C = A @ B



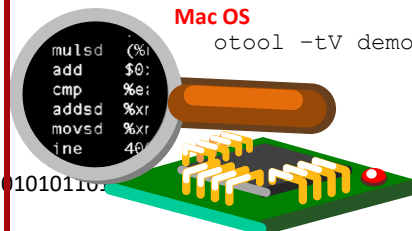
Compilation/Interpretation

Linux

```
objdump -d demo.x
```

Mac OS

```
otool -tV demo.x
```

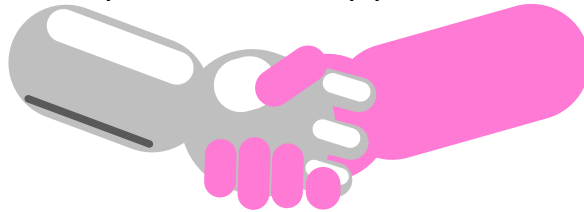


© 2025 Tze Meng Low

4

Instruction Set Architecture

- Interface between hardware and software
 - Lowest level of control for the programmer
 - Exposes program details such as
 - Program state
 - Instructions used
 - Registers
 - Relatively stable for many years / decades

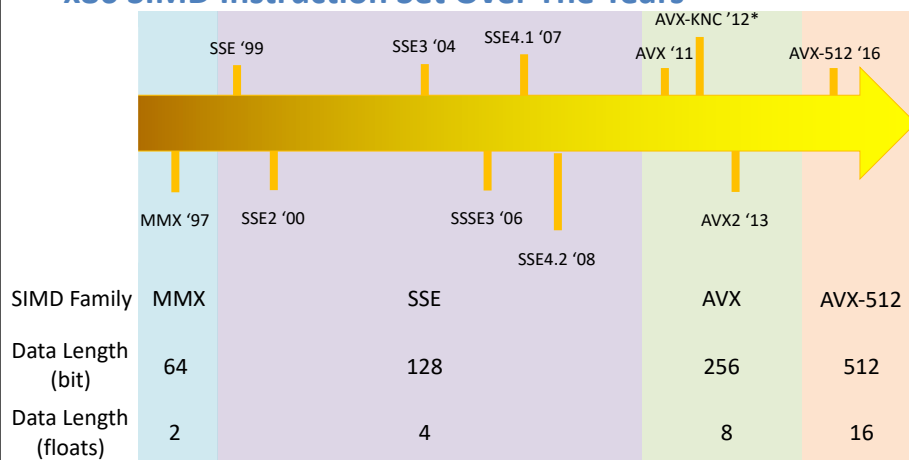


© 2025 Tze Meng Low

5

Instruction Set Architecture

- ISA can change to introduce new capabilities
- x86 SIMD Instruction Set Over The Years**



© 2025 Tze Meng Low

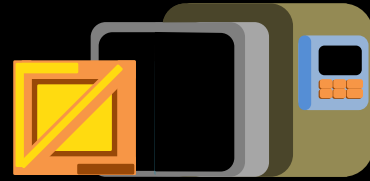
6

ISA is an abstraction

- Instructions may be executed in different order

ISA Abstraction

Behind the scene

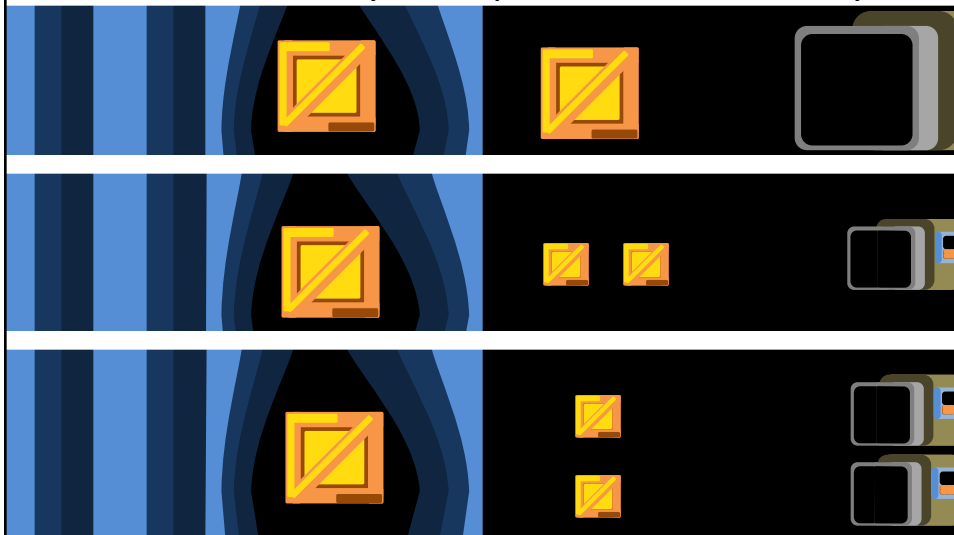


© 2025 Tze Meng Low

9

ISA is an abstraction

- Instructions may be implemented differently



10

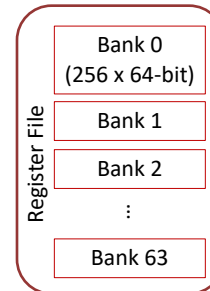
ISA is an abstraction

■ Logical registers vs physical registers

Small number of logical registers
(16 or 32)

```
mulsd    (%r8,%rax,8),%xmm0
add      $0x1,%rax
cmp      %eax,%r12d
addsd    %xmm0,%xmm1
movsd    %xmm1, (%rdx,%rdi,1)
jne      400740 <main+0x190>
```

Large physical register files
(100s to 1000s)



© 2025 Tze Meng Low

11

What about the HW do we need to know?

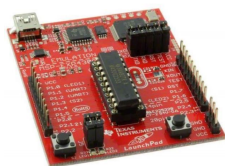
Personal Laptop/Desktop



Cloud Services



Microcontrollers



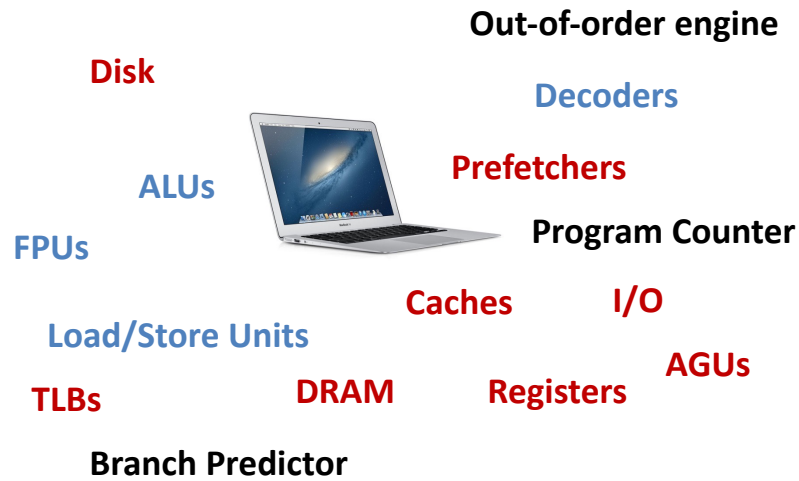
Accelerators



© 2025 Tze Meng Low

12

What about the HW do we need to know?

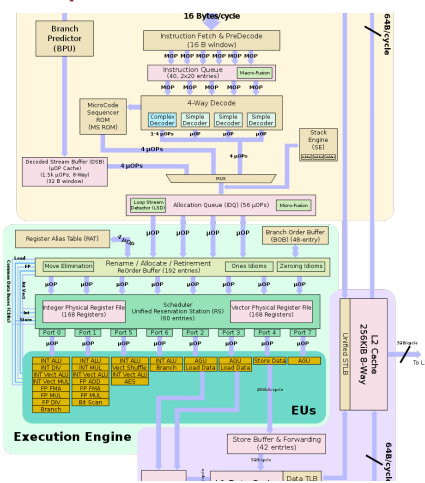


© 2025 Tze Meng Low

13

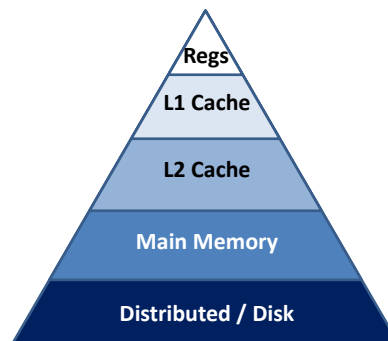
Two main parts – Execution & Data

Computation



<https://en.wikichip.org/wiki/intel/microarchitectures/haswell>

Data Movement



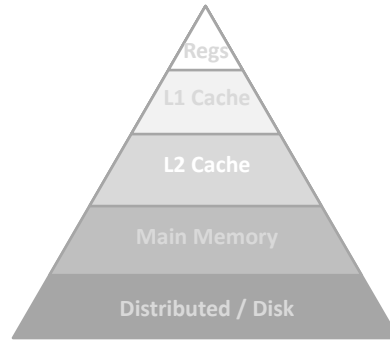
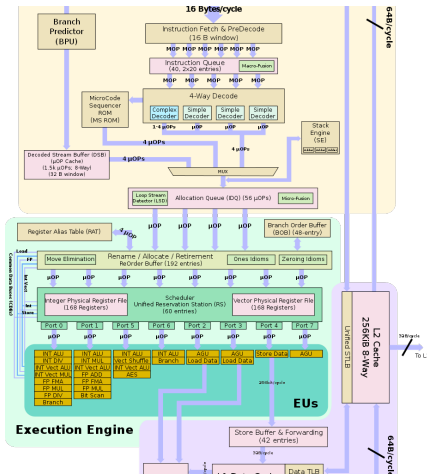
© 2025 Tze Meng Low

14

Two main parts – Execution & Data

Computation

Data Movement



<https://en.wikipedia.org/wiki/intel/microarchitectures/haswell>

© 2025 Tze Meng Low

15

A short diversion

© 2025 Tze Meng Low

16

How to optimize this?



© 2025 Tze Meng Low

17

Scenario 1



Cashier's responsibility

- Scan Items
- Bag Items
- Check receipt / Distribute free gift

On Average:

- 2 min for each task

■ Questions

- How many customer (on average) every hour?
- How long does each customer take?

© 2025 Tze Meng Low

18

You have been hired to speed up the supermarket checkout process

What are the different ways the checkout process can be improved?

- Explain why they help in speeding up the checkout process

How are they similar to hardware features we see in today's architecture?

Back to architecture

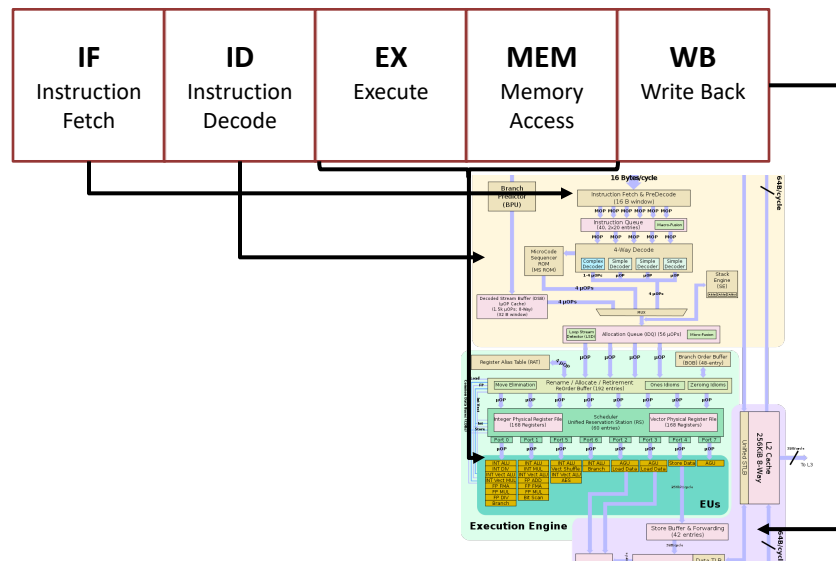
Many functional units

- Many queues/pipelines in hardware

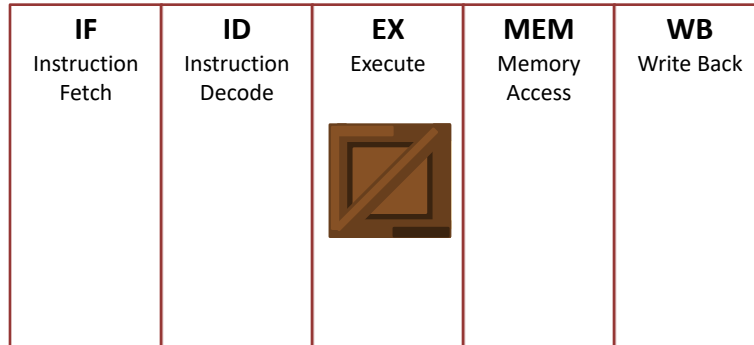
INT ALU	INT ALU	INT ALU	INT ALU	AGU	AGU	Store Data	AGU
INT DIV	INT MUL	Vect Shuffle	Branch	Load Data	Load Data		
INT Vect ALU	INT Vect ALU	INT Vect ALU					
INT Vect MUL	FP ADD	AES					
FP FMA	FP FMA						
FP MUL	FP MUL						
FP DIV	Bit Scan						
Branch							

- Goal of “Fast code” developer
 - Use as many pipelines as possible
 - Avoid stalling any of the pipelines

Model Instruction Pipeline



Instruction Pipeline



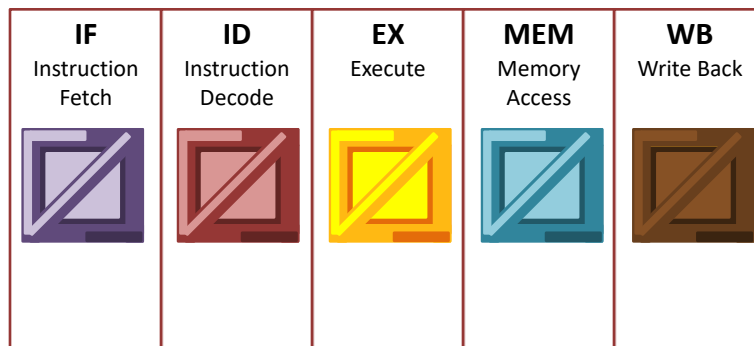
Most of the hardware pipeline are empty.
Fast code fills in the pipelines as much as possible



© 2025 Tze Meng Low

23

Instruction Pipeline



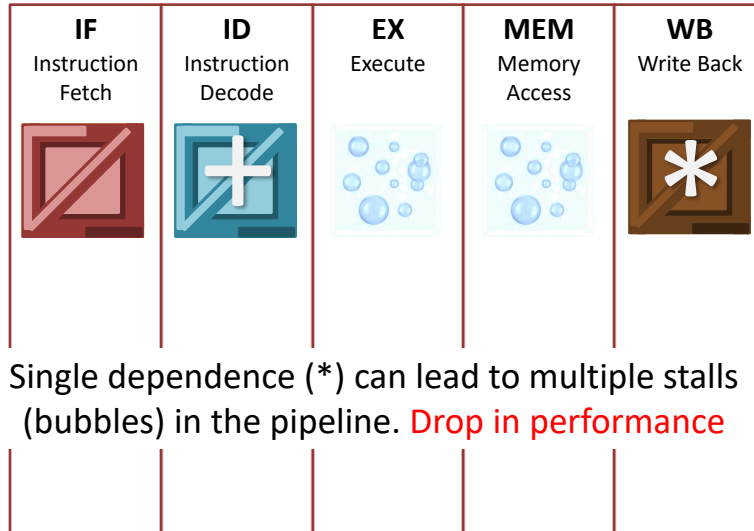
Different instructions can be in different stages



© 2025 Tze Meng Low

24

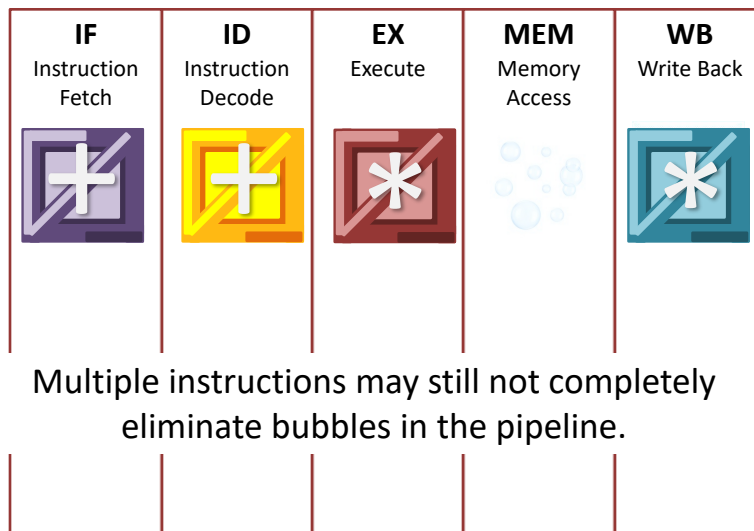
Instruction Pipeline

 $c += a * b;$


© 2025 Tze Meng Low

26

Instruction Pipeline

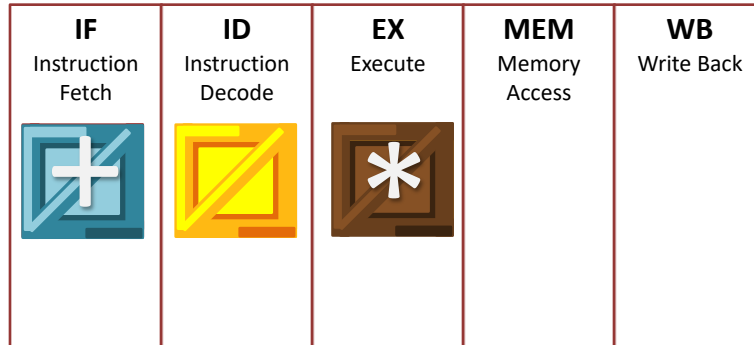
 $c += a0 * b0;$
 $c += a1 * b1;$


© 2025 Tze Meng Low

27

Instruction Pipeline

$c += a * b;$



**Need to find independent instructions!
Reorder instructions!**



© 2025 Tze Meng Low

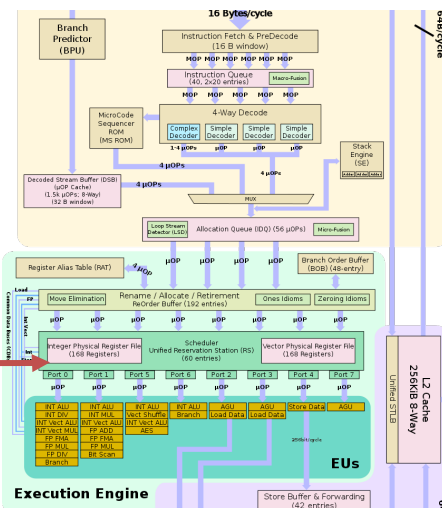
29

Instruction (Re)scheduling

Hardware solution

Out-of-order Execution

- Instruction buffer
- Fire instruction when data is available
- Write-back stage needs to be managed to ensure correctness




© 2025 Tze Meng Low

30

Instruction (Re)scheduling

- Compiler solution
 - Software pipelining
 - Compiler tries to identify and moves the instructions during the compilation process



```

mulsd    (%r8,%rax,8),%xmm0
add      $0x1,%rax
cmp      %eax,%r12d
addsd    %xmm0,%xmm1
movsd    %xmm1, (%rdx,%rdi,1)
jne      400740 <main+0x190>

```

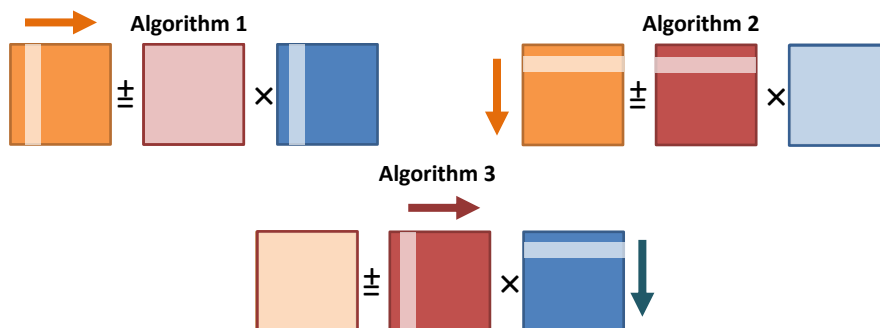
Remember obj dump ?

© 2025 Tze Meng Low

32

Recap: Matrix Multiply as an Example

$$C \pm A \times B$$



All $O(n^3)$. Which Algorithm?

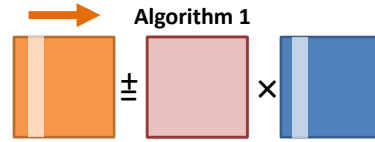
© 2025 Tze Meng Low

33

Recap: Implementations of MMM

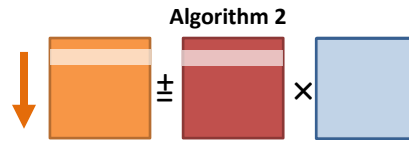
Algorithm 1

```
for (int j = 0; j != n; ++j)
  for (int i = 0; i != m; ++i)
    for (int p = 0; p != k; ++p)
      c[j*m+i] += a[p*m+i] * b[j*k+p];
```



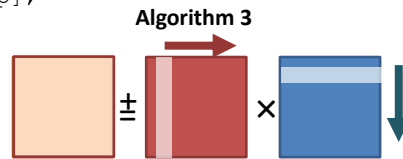
Algorithm 2

```
for (int i = 0; i != m; ++i)
  for (int j = 0; j != n; ++j)
    for (int p = 0; p != k; ++p)
      c[j*m+i] += a[p*m+i] * b[j*k+p];
```



Algorithm 3

```
for (int p = 0; p != k; ++p)
  for (int i = 0; i != m; ++i)
    for (int j = 0; j != n; ++j)
      c[j*m+i] += a[p*m+i] * b[j*k+p];
```



© 2025 Tze Meng Low

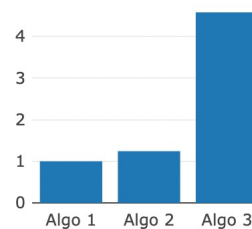
34

Matrix Multiply as an Example

- Using mmm.html

Loop Order	Execution Cycles
i, j, k	93 cycles
j, i, k	71 cycles
k, i, j	18 cycles

Higher is better



Why?

Hint: Consider what is the basic sequence of operations for each algorithm

© 2025 Tze Meng Low

35

Summary

- ISA is the “lowest” level exposed to the programmer
- ASM does not mean performance
- Instruction rescheduling is and **MUST BE** implemented at HW, SW and algorithm level
- Algorithms must be matched to hardware features