

Know Your Machine - Benchmarking

18-645: How to write fast code I

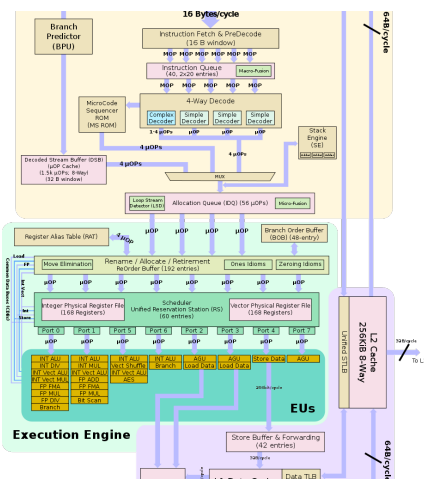
Tze Meng Low

© 2025 Tze Meng Low


1

Recap: Two main parts – Execution & Data

Computation



Data Movement



<https://en.wikichip.org/wiki/intel/microarchitectures/haswell>

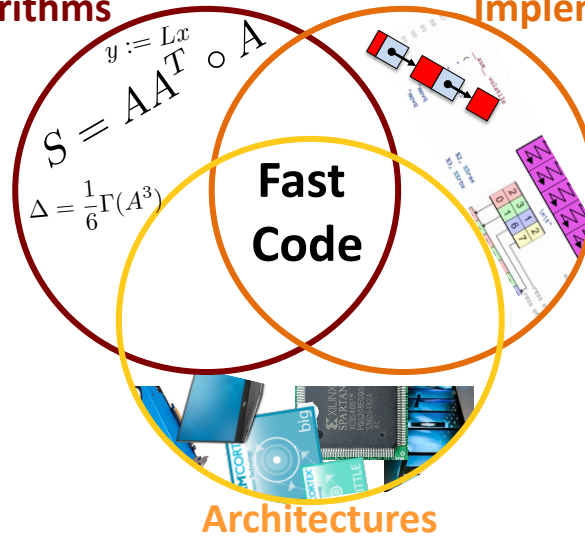
© 2025 Tze Meng Low

2

Recap: Same idea at different level

Algorithms

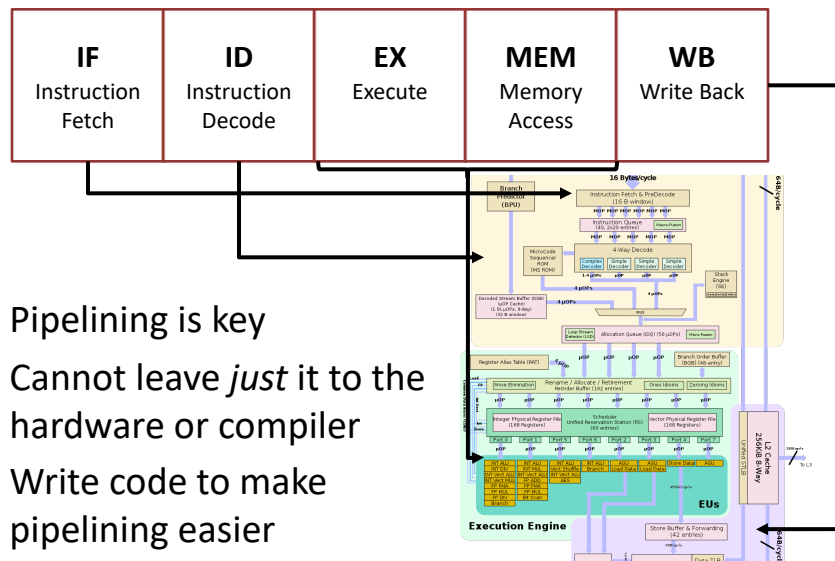
Implementations



© 2025 Tze Meng Low

3

Model Instruction Pipeline

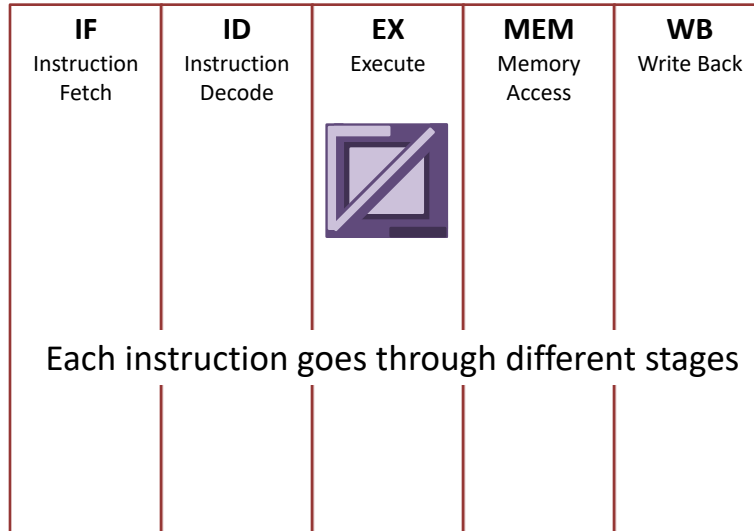


- Pipelining is key
- Cannot leave *just* it to the hardware or compiler
- Write code to make pipelining easier

© 2025 Tze Meng Low

4

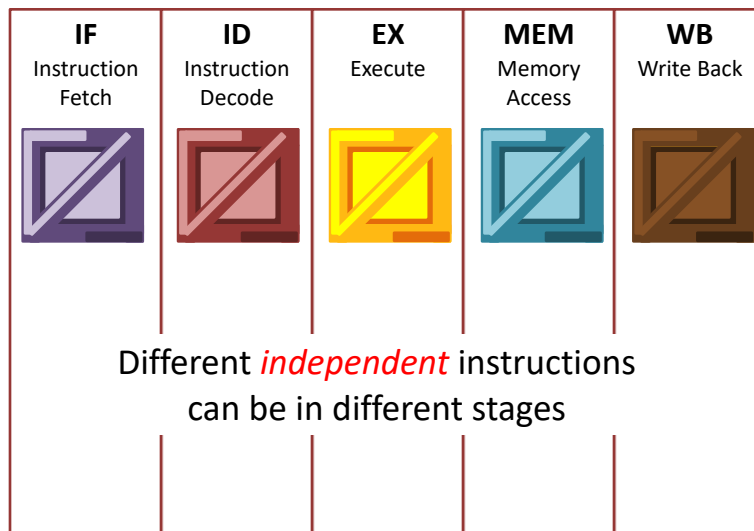
Pipelining in a nutshell



© 2025 Tze Meng Low

7

Pipelining in a nutshell



© 2025 Tze Meng Low

10

Many functional units

- Many specialized queues/pipelines in HW

INT ALU	INT ALU	INT ALU	INT ALU	AGU	AGU	Store Data	AGU
INT DIV	INT MUL	Vect Shuffle	Branch	Load Data	Load Data		
INT Vect ALU	INT Vect ALU	INT Vect ALU					
INT Vect MUL	FP ADD	AES					
FP FMA	FP FMA						
FP MUL	FP MUL						
FP DIV	Bit Scan						
Branch							

- Goal of “Fast code” developer
 - Use as many pipelines as possible
 - Avoid stalling any of the pipelines
 - Find independent instructions
- How many instructions do we need?**

© 2025 Tze Meng Low

12

Benchmarking

- Want to know determine machine capability
 - Latency of instructions
 - Minimum # cycles before next *dependent* instruction
 - Instruction throughput
 - Instructions processed per *clock cycle*
- Benchmarking is empirical
 - It will be messy and time-consuming (*Start HW early*)
 - Do not expect perfect answers
 - Getting accurate timings is difficult
- Key assumption: Instruction is fully pipelined**



© 2025 Tze Meng Low

13

Timing

■ Methods for timing

— x86 (single-core performance)

- Wall clock time `gettimeofday`
- `rdtsc` (returns time stamp counter) `__rdtsc`

```
unsigned long long rdtsc()
{
    unsigned long long int x;
    unsigned a, d;

    __asm__ volatile("rdtsc" : "=a" (a), "=d" (d));

    return ((unsigned long long)a) | (((unsigned long long)d) << 32);
}
```

- returns **clock cycles** at nominal frequency

— Other architectures

- Clock events on GPU (milliseconds)

© 2025 Tze Meng Low

14

Previously

IF	ID	EX	MEM	WB
Instruction Fetch	Instruction Decode	Execute	Memory Access	Write Back


- Assumed that each stage take 1 time unit
 - 1 cycle or 1 clock tick
- Total of 5 independent instructions can be “in flight” or executed

© 2025 Tze Meng Low

15

On more realistic architectures

- Details of pipeline is typically unknown
- Instead, we want to track how long before an instruction is “done”



© 2025 Tze Meng Low

17

Latency

- Minimum # cycles before next dependent instruction

▪ Process:

```
st = rdtsc();
int var = 0;
for (int i = 0; i != REPS; i++)
    var++;
et = rdtsc();
```

- Metric

$$\frac{\text{Time in Cycles}}{\text{Number of Instructions}} = \frac{et - st}{REPS}$$

What is/are the problem(s)?

© 2025 Tze Meng Low

20

Multiple instructions for the same code

■ Difficulties:

- Compiler may not select desired instructions

```
float b = 1.0;
float a = 0.0;
a += b;
a += b;
...
```

addsd
fadd

Both are floating point addition.

CHECK YOUR ASSEMBLY!!

- Compiler may optimize instructions out

```
float b = 1.0;
float a = 0.0;
a += b;
a += b;
...
a += b;
```

Compiler optimization

```
float b = 1.0;
float a = 100.0;
```

Mac OS

otool -tv demo.x

Linux

objdump -d demo.x

© 2025 Tze Meng Low

24

Macro Intrinsics

- Use C macros that hide inline assembly
- Write in macros instead

```
#define MULTIPLY(dest, src)
__asm__ __volatile__(
    "imul %[rsrc], %[rdest]\n"
    : [rdest] "+r" (dest)
    : [rsrc] "r" (src)
    );
```

Try not move code

Assembly to test

Input is read from register

Destination is modified (+)
Destination is a register (r)

Other operand descriptors

m - memory

x - SIMD register

Veras et al, Compilers, hands-off my hands-on optimizations, Proceedings of the 3rd Workshop on Programming Models for SIMD/Vector Processing, March 2016

© 2025 Tze Meng Low

28

Macro Intrinsics

- How to use

As a single instruction

```
int x = 1;
int y = 2;

//x *= y
MULTIPLY(x, y)
```

As part of macro

```
#define MULTIPLY3(x, a) \
    MULTIPLY(x, a)      \
    MULTIPLY(x, a)      \
    MULTIPLY(x, a)

int x = 1;
int y = 2;

// x = ((x * y) * y) * y
MULTIPLY3(x, y)
```

© 2025 Tze Meng Low

29

Latency

- Minimum # cycles before next dependent instruction

Process:

```
int var = 0;
int tmp = 1;
st = rdtsc();
ADD100(var, tmp); //assume REPS == 100
et = rdtsc();
```

- Metric

$$\frac{\text{Time in Cycles}}{\text{Number of Instructions}} = \frac{et - st}{REPS}$$

© 2025 Tze Meng Low

30

Latency

- Minimum # cycles before next dependent instruction
- Process:
 - Time a chain of dependent instructions
 - Divide time by number of instructions in chain
- Notes:
 - `rdtsc` may not be accurate for small cycle counts
 - Create long (> 100s – 1000s) chains
 - ***This is empirical. Don't expect exact numbers. If in doubt, round up***

© 2025 Tze Meng Low

31

Throughput

- Instructions processed per unit time
- Notice that some of the instructions can be processed by multiple functional units

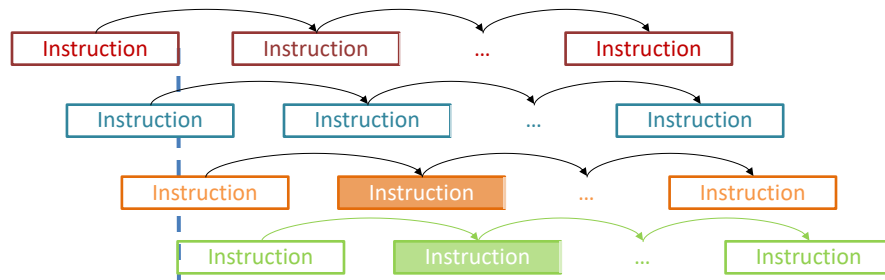
INT ALU	INT ALU	INT ALU	INT ALU	AGU	AGU	Store Data	AGU
INT DIV	INT MUL	Vect Shuffle	Branch	Load Data	Load Data		
INT Vect ALU	INT Vect ALU	INT Vect ALU					
INT Vect MUL	FP ADD	AES					
FP FMA	FP FMA						
FP MUL	FP MUL						
FP DIV	Bit Scan						
Branch							

© 2025 Tze Meng Low

32

Throughput

- Assume only 1 functional unit, 3 cycles latency



How many longer does it takes when another
Execution Time **independent chain** of **dependent instructions**?



© 2025 Tze Meng Low

47

Throughput benchmarking

- Process:
 - Time 1 chain of dependent instructions
 - Time increasing **interleaved** chains of dependent instructions
 - Find largest number of chains **before** execution time increases significantly
 - Compute throughput

$$\text{Throughput} = \frac{\#Instr. Per Chain \times \#Chains}{Time in Cycles}$$

© 2025 Tze Meng Low

48


Interleaving code is necessary

- HW has limited capability to interleave code
- Manual interleaving (use Macro intrinsics)

```

c += a[0] * b[0]
c += a[1] * b[1]
c += a[2] * b[2]
:
c += a[998] * b[998]
c += a[999] * b[999]
c += a[1000] * b[1000]
d += a[0] * b[0]
d += a[1] * b[1]
d += a[2] * b[2]
:
d += a[998] * b[998]
d += a[999] * b[999]
d += a[1000] * b[1000]



```



```

c += a[0] * b[0]
d += a[1] * b[1]
c += a[2] * b[2]
d += a[2] * b[2]
:
c += a[998] * b[998]
d += a[998] * b[998]
c += a[999] * b[999]
d += a[999] * b[999]
c += a[1000] * b[1000]
d += a[1000] * b[1000]
:

```



© 2025 Tze Meng Low

49

Throughput

- Instructions processed per unit time

2 customers leave queues **every 6 minutes**1 customer leave queue **every 3 minutes**

1 customers / 3 min

© 2025 Tze Meng Low

50

Hardware information on the web

Code Name	Products formerly Coffee Lake
Vertical Segment	Desktop
Processor Number	i9-9900T
Status	Launched
Launch Date ?	Q2'19
Lithography ?	14 nm
Use Conditions ?	PC/Client/Tablet
Recommended Customer Price ?	\$439.00

Performance

# of Cores ?	8
# of Threads ?	16
Processor Base Frequency ?	2.10 GHz
Max Turbo Frequency ?	4.40 GHz
Cache ?	16 MB SmartCache
Bus Speed ?	8 GT/s DMI3
TDP ?	35 W

<https://ark.intel.com/content/www/us/en/ark/products/191044/intel-core-i9-9900t-processor-16m-cache-up-to-4-40-ghz.html>

© 2025 Tze Meng Low

51

Hardware information on the web

Floating point XMM and YMM instructions

Instruction	Operands	µops fused domain	µops unfused domain						Latency	Reciprocal throughput	Comments
			p015	p0	p1	p5	p23	p4			
Move instructions											
MOVAPS/D	x,x	1	1			1			0-1	≤1	elimin.
VMOVAPS/D	y,y	1	1			1			0-1	≤1	elimin.
MOVAPS/D MOVUPS/D	x,m128	1					1		3	0.5	
VMOVAPS/D											
VMOVUPS/D	y,m256	1					1+		4	1	AVX
MOVAPS/D MOVUPS/D	m128,x	1					1	1	3	1	

WARNING! Many such sites, not all are accurate!
Same information, Reported differently
Similar but not identical architectures

<https://uops.info/>

https://www.agner.org/optimize/instruction_tables.pdf

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

© 2025 Tze Meng Low

52

Summary

- Benchmarking is important for knowing information we need for designing fast code
- Benchmarking is messy and is not precise
- These information can be found through the writing of small programs (micro-benchmark)